

Data Science for Economists

Lecture 5: Data cleaning & wrangling: (1) Tidyverse

Grant McDermott, adapted by Kyle Coombs

Bates College | [EC/DCS 368](#)

Table of contents

1. Prologue
2. Tidyverse basics
3. Data wrangling with dplyr
 - filter
 - arrange
 - select
 - mutate
 - summarise
 - joins
4. Data tidying with tidyr
 - pivot_longer / pivot_wider
5. Summary

Prologue

Why so many packages?

- You are probably wondering why there are so many packages in R that do similar things.
- How come you need to know this many packages? Isn't this a bit much?
- Think back to our clean code principles.
 - One of the key practices of clean code is to abstract away complexity.
 - This is what packages do. They abstract away the complexity to make code easier to read, write, and debug.
 - They offer a consistent interface and set of help documentation.
 - Different packages prioritize different goals -- so you can choose the one that best fits your needs.
 - e.g. the `tidyverse` packages prioritize relational database management (called "tidy" data)
 - `data.table` prioritizes speed and memory efficiency in completing data operations, assumes you're doing the RDBM yourself

Why so many packages?

- You are probably wondering why there are so many packages in R that do similar things.
- How come you need to know this many packages? Isn't this a bit much?
- Think back to our clean code principles.
 - One of the key practices of clean code is to abstract away complexity.
 - This is what packages do. They abstract away the complexity to make code easier to read, write, and debug.
 - They offer a consistent interface and set of help documentation.
 - Different packages prioritize different goals -- so you can choose the one that best fits your needs.
 - e.g. the `tidyverse` packages prioritize relational database management (called "tidy" data)
 - `data.table` prioritizes speed and memory efficiency in completing data operations, assumes you're doing the RDBM yourself
- Of course, different packages have different ways of abstracting away complexity.
- So yes, it is a bit much, but it's also a good thing.

Checklist

R packages you'll need for this lecture

☑ **tidyverse**

- This is a meta-package that loads a suite of other packages, including **dplyr** and **tidyr**, which includes the `starwars` dataset that we'll use for practice.

☑ **nycflights13**

Checklist

R packages you'll need for this lecture

☑ **tidyverse**

- This is a meta-package that loads a suite of other packages, including **dplyr** and **tidyr**, which includes the **starwars** dataset that we'll use for practice.

☑ **nycflights13**

The following code chunk will install (if necessary) and load everything for you.

```
if (!require(pacman)) install.packages('pacman', repos = 'https://cran.rstudio.com')
pacman::p_load(tidyverse, nycflights13)
```

What is "tidy" data?

Resources:

- [Vignettes](#) (from the **tidyr** package)
- [Original paper](#) (Hadley Wickham, 2014 JSS)

What is "tidy" data?

Resources:

- [Vignettes](#) (from the **tidyr** package)
- [Original paper](#) (Hadley Wickham, 2014 JSS)

Key points:

1. Each variable forms a column.
2. Each observation forms a row.
3. Each type of observational unit forms a table.

What is "tidy" data?

Resources:

- [Vignettes](#) (from the **tidyr** package)
- [Original paper](#) (Hadley Wickham, 2014 JSS)

Key points:

1. Each variable forms a column.
2. Each observation forms a row.
3. Each type of observational unit forms a table.

Basically, tidy data is more likely to be [long \(i.e. narrow\) format](#) than wide format.

Relational Database Management with R

- Remember Relational Database Management from our work on [Empirical Organization?](#)
- Today, we'll learn how to implement it using packages in the tidyverse
- We'll cover:
 - Subsetting data
 - Variable creation, renaming, selection
 - Grouping and summarizing data
 - Joining and appending datasets

Tidyverse basics

Tidyverse vs. base R

There is often a direct correspondence between a **tidyverse** command and its **base R** equivalent.

These generally follow a `tidyverse::snake_case` vs `base::period.case` rule:

tidyverse	base
<code>?readr::read_csv</code>	<code>?utils::read.csv</code>
<code>?dplyr::if_else</code>	<code>?base::ifelse</code>
<code>?tibble::tibble</code>	<code>?base::data.frame</code>

Etcetera.

If you call up the above examples, you'll see that the tidyverse alternative:

- Offers enhancements or other useful options (and some restrictions too)
- Better documentation
- More consistent syntax

Tidyverse vs. base R

There is often a direct correspondence between a **tidyverse** command and its **base R** equivalent.

These generally follow a `tidyverse::snake_case` vs `base::period.case` rule:

tidyverse	base
<code>?readr::read_csv</code>	<code>?utils::read.csv</code>
<code>?dplyr::if_else</code>	<code>?base::ifelse</code>
<code>?tibble::tibble</code>	<code>?base::data.frame</code>

Etcetera.

If you call up the above examples, you'll see that the tidyverse alternative:

- Offers enhancements or other useful options (and some restrictions too)
- Better documentation
- More consistent syntax

Remember: There are (almost) always multiple ways to achieve a single goal in R.

Tidyverse packages

Let's load the tidyverse meta-package and check the output.

```
library(tidyverse)
```

Tidyverse packages

Let's load the tidyverse meta-package and check the output.

```
library(tidyverse)
```

We have actually loaded a number of packages (which could also be loaded individually):

ggplot2, **tibble**, **dplyr**, etc.

- We can also see information about the package versions and some [namespace conflicts](#).

Tidyverse packages (cont.)

The tidyverse actually comes with a lot more packages than those loaded automatically.¹

```
tidyverse_packages()
```

```
## [1] "broom"          "conflicted"    "cli"           "dbplyr"
## [5] "dplyr"          "dtplyr"        "forcats"       "ggplot2"
## [9] "googledrive"    "googlesheets4" "haven"         "hms"
## [13] "httr"           "jsonlite"      "lubridate"     "magrittr"
## [17] "modelr"         "pillar"        "purrr"         "ragg"
## [21] "readr"          "readxl"        "reprex"        "rlang"
## [25] "rstudioapi"     "rvest"         "stringr"       "tibble"
## [29] "tidyr"          "xml2"          "tidyverse"
```

We'll use most of these packages during the remainder of this course.

- **lubridate** for dates, **rvest** for webscraping, **broom** to `tidy()` R objects into tables
- However, packages still have to be loaded separately with `library()`

¹ It also includes a *lot* of dependencies upon installation. This is a matter of some [controversy](#).

Tidyverse packages (cont.)

Today, however, I'm only really going to focus on two packages:

1. **dplyr**
2. **tidyr**

These are the workhorse packages for cleaning and wrangling data.

- Data cleaning and wrangling occupies an inordinate amount of time, no matter where you are in your research career.
- I cannot underscore this enough
- This course can add structure to the cleaning and wrangling, but it is still a time-consuming process.
- It can be a real bummer, so pick data projects that you are excited about.

dplyr

Key dplyr verbs

There are five key dplyr verbs that you need to learn.

1. `filter`: Filter (i.e. subset) rows based on their values.
2. `arrange`: Arrange (i.e. reorder) rows based on their values.
3. `select`: Select (i.e. subset) columns by their names:
4. `mutate`: Create new columns.
5. `summarise`: Collapse multiple rows into a single summary value.¹

¹ `summarize` with a "z" works too, but Hadley Wickham is from New Zealand.

Learn the verbs

Practice these commands together using the `starwars` data frame that comes pre-packaged with dplyr. **Stop** when you hit the last `summarise` slide (approx. 33).

```
starwars
```

```
## # A tibble: 87 × 14
##   name      height  mass hair_color skin_color eye_color birth_year sex  gender
##   <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
## 1 Luke Sk...   172    77 blond      fair        blue         19  male  mascu...
## 2 C-3PO        167    75 <NA>      gold        yellow       112  none  mascu...
## 3 R2-D2         96    32 <NA>      white, bl... red          33  none  mascu...
## 4 Darth V...   202   136 none       white       yellow      41.9  male  mascu...
## 5 Leia Or...   150    49 brown      light       brown        19  fema... femin...
## 6 Owen La...   178   120 brown, gr... light       blue         52  male  mascu...
## 7 Beru Wh...   165    75 brown      light       blue         47  fema... femin...
## 8 R5-D4         97    32 <NA>      white, red  red          NA  none  mascu...
## 9 Biggs D...   183    84 black      light       brown        24  male  mascu...
## 10 Obi-Wan...   182    77 auburn, w... fair        blue-gray    57  male  mascu...
## # i 77 more rows
## # i 5 more variables: homeworld <chr>, species <chr>, films <list>,
## #   vehicles <list>, starships <list>
```

1) dplyr::filter

Filter means "subset" the rows of a data frame based on some condition(s).

```
starwars %>%  
  filter(species == "Human", height ≥ 190)
```

```
## # A tibble: 4 × 14  
##   name      height  mass hair_color skin_color eye_color birth_year sex  gender  
##   <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>  
## 1 Darth Va...   202   136 none      white      yellow      41.9 male  masculi  
## 2 Qui-Gon ...   193    89 brown     fair       blue        92  male  masculi  
## 3 Dooku        193    80 white     fair       brown       102  male  masculi  
## 4 Bail Pre...   191    NA black     tan        brown        67  male  masculi  
## # i 5 more variables: homeworld <chr>, species <chr>, films <list>,  
## #   vehicles <list>, starships <list>
```

We can chain multiple commands with the pipe `%>%` as we've seen¹.

¹ Pipes were invented by Doug McIlroy in 1964, are widely used in Unix shells (e.g. bash) and other programming languages (e.g. F#). They pass the preceding object as the first argument to the following function. In R, they allow you to chain together code in a way that reads from left to right.

1) dplyr::filter cont.

A very common `filter` use case is identifying (or removing) missing data cases.

```
starwars %>%  
  filter(is.na(height))
```

```
## # A tibble: 6 × 14  
##   name      height  mass hair_color skin_color eye_color birth_year sex  gender  
##   <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>  
## 1 Arvel Cr...    NA     NA brown      fair       brown             NA male  mascu...  
## 2 Finn          NA     NA black      dark       dark             NA male  mascu...  
## 3 Rey           NA     NA brown      light      hazel            NA fema... femin...  
## 4 Poe Dame...    NA     NA brown      light      brown            NA male  mascu...  
## 5 BB8           NA     NA none       none       black            NA none  mascu...  
## 6 Captain ...    NA     NA unknown   unknown    unknown          NA <NA>   <NA>  
## # i 5 more variables: homeworld <chr>, species <chr>, films <list>,  
## #   vehicles <list>, starships <list>
```

To remove missing observations, simply use negation: `filter(!is.na(height))`. Try this yourself.

2) dplyr::arrange

```
starwars %>%  
  arrange(birth_year)
```

```
## # A tibble: 87 × 14  
##   name      height  mass hair_color skin_color eye_color birth_year sex  gender  
##   <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>  
## 1 Wicket ...    88  20  brown      brown      brown         8  male  masculi  
## 2 IG-88        200 140  none       metal      red          15  none  masculi  
## 3 Luke Sk...   172  77  blond      fair       blue         19  male  masculi  
## 4 Leia Or...   150  49  brown      light      brown         19  fema... femin...  
## 5 Wedge A...   170  77  brown      fair       hazel         21  male  masculi  
## 6 Plo Koon     188  80  none       orange     black         22  male  masculi  
## 7 Biggs D...   183  84  black      light      brown         24  male  masculi  
## 8 Han Solo     180  80  brown      fair       brown         29  male  masculi  
## 9 Lando C...   177  79  black      dark       brown         31  male  masculi  
## 10 Boba Fe...   183  78.2 black      fair       brown        31.5  male  masculi  
## # i 77 more rows  
## # i 5 more variables: homeworld <chr>, species <chr>, films <list>,  
## #   vehicles <list>, starships <list>
```


2) dplyr::arrange

```
starwars %>%  
  arrange(birth_year)
```

```
## # A tibble: 87 × 14  
##   name      height  mass hair_color skin_color eye_color birth_year sex  gender  
##   <chr>    <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>  
## 1 Wicket ...    88  20  brown     brown     brown         8  male  mascu...  
## 2 IG-88        200 140  none      metal      red          15  none  mascu...  
## 3 Luke Sk...   172  77  blond     fair       blue         19  male  mascu...  
## 4 Leia Or...   150  49  brown     light      brown        19  fema... femin...  
## 5 Wedge A...   170  77  brown     fair       hazel        21  male  mascu...  
## 6 Plo Koon    188  80  none      orange     black        22  male  mascu...  
## 7 Biggs D...   183  84  black     light      brown        24  male  mascu...  
## 8 Han Solo    180  80  brown     fair       brown        29  male  mascu...  
## 9 Lando C...   177  79  black     dark       brown        31  male  mascu...  
## 10 Boba Fe...  183  78.2 black     fair       brown       31.5  male  mascu...  
## # i 77 more rows  
## # i 5 more variables: homeworld <chr>, species <chr>, films <list>,  
## #   vehicles <list>, starships <list>
```

Note: Arranging on a character-based column (i.e. strings) will sort alphabetically. Try this yourself by arranging according to the "name" column.

2) dplyr::arrange cont.

We can also arrange items in descending order using `arrange(desc())`.

```
starwars %>%  
  arrange(desc(birth_year))
```

```
## # A tibble: 87 × 14  
##   name      height  mass hair_color skin_color eye_color birth_year sex  gender  
##   <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>  
## 1 Yoda         66    17 white      green      brown         896 male  mascu...  
## 2 Jabba D...   175   1358 <NA>      green-tan... orange        600 herm... mascu...  
## 3 Chewbac...   228   112 brown     unknown    blue         200 male  mascu...  
## 4 C-3PO       167    75 <NA>      gold       yellow        112 none  mascu...  
## 5 Dooku        193    80 white     fair       brown        102 male  mascu...  
## 6 Qui-Gon...   193    89 brown     fair       blue          92 male  mascu...  
## 7 Ki-Adi-...   198    82 white     pale       yellow         92 male  mascu...  
## 8 Finis V...   170    NA blond     fair       blue          91 male  mascu...  
## 9 Palpati...   170    75 grey     pale       yellow         82 male  mascu...  
## 10 Cliegg ...   183    NA brown     fair       blue          82 male  mascu...  
## # i 77 more rows  
## # i 5 more variables: homeworld <chr>, species <chr>, films <list>,  
## #   vehicles <list>, starships <list>
```

3) dplyr::select

Select means subset the columns of a data frame based on their names.

Use commas to select multiple columns out of a data frame. (You can also use "first:last" for consecutive columns). Deselect a column with "-".

```
starwars %>%  
  select(name:skin_color, species, -height) %>%  
  head()
```

```
## # A tibble: 6 × 5  
##   name          mass hair_color skin_color species  
##   <chr>        <dbl> <chr>      <chr>      <chr>  
## 1 Luke Skywalker    77 blond      fair        Human  
## 2 C-3PO             75 <NA>      gold        Droid  
## 3 R2-D2             32 <NA>      white, blue Droid  
## 4 Darth Vader      136 none       white        Human  
## 5 Leia Organa       49 brown      light        Human  
## 6 Owen Lars        120 brown, grey light        Human
```

3) dplyr::select *cont.*

You can also rename some (or all) of your selected variables in place.

```
starwars %>%  
  select(alias=name, crib=homeworld, sex=gender) %>%  
  head()
```

```
## # A tibble: 6 × 3  
##   alias      crib      sex  
##   <chr>      <chr>    <chr>  
## 1 Luke Skywalker Tatooine masculine  
## 2 C-3PO      Tatooine masculine  
## 3 R2-D2      Naboo     masculine  
## 4 Darth Vader Tatooine masculine  
## 5 Leia Organa Alderaan feminine  
## 6 Owen Lars  Tatooine masculine
```

3) dplyr::select *cont.*

You can also rename some (or all) of your selected variables in place.

```
starwars %>%  
  select(alias=name, crib=homeworld, sex=gender) %>%  
  head()
```

```
## # A tibble: 6 × 3  
##   alias      crib      sex  
##   <chr>      <chr>    <chr>  
## 1 Luke Skywalker Tatooine masculine  
## 2 C-3PO      Tatooine masculine  
## 3 R2-D2      Naboo     masculine  
## 4 Darth Vader Tatooine masculine  
## 5 Leia Organa Alderaan feminine  
## 6 Owen Lars  Tatooine masculine
```

If you just want to rename columns without subsetting them, you can use `rename`. Try this now by replacing `select(...)` in the above code chunk with `rename(...)`.

4) dplyr::mutate

You can create new columns from scratch, or (more commonly) as transformations of existing columns.

```
starwars %>%  
  select(name, birth_year) %>%  
  mutate(dog_years = birth_year * 7) %>%  
  mutate(comment = paste0(name, " is ", dog_years, " in dog years.")) %>%  
  head()
```

```
## # A tibble: 6 × 4  
##   name          birth_year dog_years comment  
##   <chr>          <dbl>     <dbl> <chr>  
## 1 Luke Skywalker      19        133 Luke Skywalker is 133 in dog years.  
## 2 C-3PO              112        784 C-3PO is 784 in dog years.  
## 3 R2-D2              33        231 R2-D2 is 231 in dog years.  
## 4 Darth Vader       41.9       293.3 Darth Vader is 293.3 in dog years.  
## 5 Leia Organa        19        133 Leia Organa is 133 in dog years.  
## 6 Owen Lars          52        364 Owen Lars is 364 in dog years.
```

4) dplyr::mutate cont.

Boolean, logical and conditional operators all work well with `mutate` too.

```
starwars %>%  
  select(name, height) %>%  
  filter(name %in% c("Luke Skywalker", "Anakin Skywalker")) %>%  
  mutate(tall1 = height > 180) %>%  
  mutate(tall2 = ifelse(height > 180, "Tall", "Short")) ## Same effect, but can choose labe
```

```
## # A tibble: 2 × 4  
##   name          height tall1 tall2  
##   <chr>         <int> <lgl> <chr>  
## 1 Luke Skywalker    172 FALSE Short  
## 2 Anakin Skywalker    188 TRUE  Tall
```

4) dplyr::mutate *cont.*

Lastly, combining `mutate` with the `across` feature allows you to easily work on a subset of variables. For example:

```
starwars %>%  
  select(name:eye_color) %>%  
  mutate(across(where(is.character), toupper)) %>%  
  head(5)
```

```
## # A tibble: 5 × 6  
##   name          height  mass hair_color skin_color eye_color  
##   <chr>         <int> <dbl> <chr>      <chr>      <chr>  
## 1 LUKE SKYWALKER    172    77 BLOND      FAIR        BLUE  
## 2 C-3PO             167    75 <NA>      GOLD        YELLOW  
## 3 R2-D2             96    32 <NA>      WHITE, BLUE RED  
## 4 DARTH VADER      202   136 NONE      WHITE        YELLOW  
## 5 LEIA ORGANA      150    49 BROWN     LIGHT        BROWN
```


5) dplyr::summarise

Particularly useful in combination with the `group_by`¹ command.

```
starwars %>%  
  group_by(species, gender) %>%  
  summarise(mean_height = mean(height, na.rm = TRUE)) %>%  
  head()
```

```
## # A tibble: 6 × 3  
## # Groups:   species [6]  
##   species gender    mean_height  
##   <chr>    <chr>         <dbl>  
## 1 Aleena   masculine         79  
## 2 Besalisk masculine        198  
## 3 Cerean   masculine        198  
## 4 Chagrian masculine        196  
## 5 Clawdite feminine        168  
## 6 Droid    feminine         96
```

Note: **dplyr** 1.0.0 also notifies you about grouping variables every time you do operations on or with them. YMMV, but I switch them off with `options(dplyr.summarise.inform = FALSE)` in my `.Rprofile`.

5) dplyr::summarise *cont.*

Note that including "na.rm = TRUE" (or, its alias "na.rm = T") is usually a good idea with summarise functions. Otherwise, your output will be missing too.

```
## Probably not what we want
starwars %>%
  summarise(mean_height = mean(height))
```

```
## # A tibble: 1 × 1
##   mean_height
##         <dbl>
## 1          NA
```

```
## Much better
starwars %>%
  summarise(mean_height = mean(height, na.rm = TRUE))
```

```
## # A tibble: 1 × 1
##   mean_height
##         <dbl>
## 1       174.
```

5) dplyr::summarise *cont.*

The same `across`-based workflow that we saw with `mutate` a few slides back also works with `summarise`. For example:

```
starwars %>%  
  group_by(species) %>%  
  summarise(across(where(is.numeric), ~mean(.x, na.rm=T))) %>%  
  head()
```

```
## # A tibble: 6 × 4  
##   species height mass birth_year  
##   <chr>    <dbl> <dbl>    <dbl>  
## 1 Aleena      79    15      NaN  
## 2 Besalisk   198   102      NaN  
## 3 Cerean     198    82      92  
## 4 Chagrian   196   NaN      NaN  
## 5 Clawdite   168    55      NaN  
## 6 Droid     131.   69.8    53.3
```

5) dplyr::summarise *cont.*

The same `across`-based workflow that we saw with `mutate` a few slides back also works with `summarise`. For example:

```
starwars %>%  
  group_by(species) %>%  
  summarise(across(where(is.numeric), ~mean(.x, na.rm=T))) %>%  
  head()
```

```
## # A tibble: 6 × 4  
##   species height mass birth_year  
##   <chr>    <dbl> <dbl>    <dbl>  
## 1 Aleena      79    15      NaN  
## 2 Besalisk   198   102      NaN  
## 3 Cerean     198    82      92  
## 4 Chagrian   196   NaN      NaN  
## 5 Clawdite   168    55      NaN  
## 6 Droid     131.   69.8    53.3
```

Try to intuit what `.x` does above!

Other dplyr goodies

`group_by` and `ungroup`: For (un)grouping.

- Particularly useful with the `summarise` and `mutate` commands, as we've already seen.

Other dplyr goodies

`group_by` and `ungroup`: For (un)grouping.

- Particularly useful with the `summarise` and `mutate` commands, as we've already seen.

`slice`: Subset rows by position rather than filtering by values.

- `starwars %>% slice(c(1, 5))`

Other dplyr goodies

`group_by` and `ungroup`: For (un)grouping.

- Particularly useful with the `summarise` and `mutate` commands, as we've already seen.

`slice`: Subset rows by position rather than filtering by values.

- `starwars %>% slice(c(1, 5))`

`pull`: Extract a column as a vector or scalar.

- `starwars %>% filter(gender="female") %>% pull(height)` returns `height` as a vector

Other dplyr goodies

`group_by` and `ungroup`: For (un)grouping.

- Particularly useful with the `summarise` and `mutate` commands, as we've already seen.

`slice`: Subset rows by position rather than filtering by values.

- `starwars %>% slice(c(1, 5))`

`pull`: Extract a column as a vector or scalar.

- `starwars %>% filter(gender="female") %>% pull(height)` returns `height` as a vector

`count` and `distinct`: Number and isolate unique observations.

- `starwars %>% count(species)`, or `starwars %>% distinct(species)`
- Or use `mutate`, `group_by`, and `n()`, e.g. `starwars %>% group_by(species) %>% mutate(num = n())`.

There are also **window functions** for leads and lags, ranks, cumulative aggregation, etc.

- See `vignette("window-functions")`.

Quick quiz

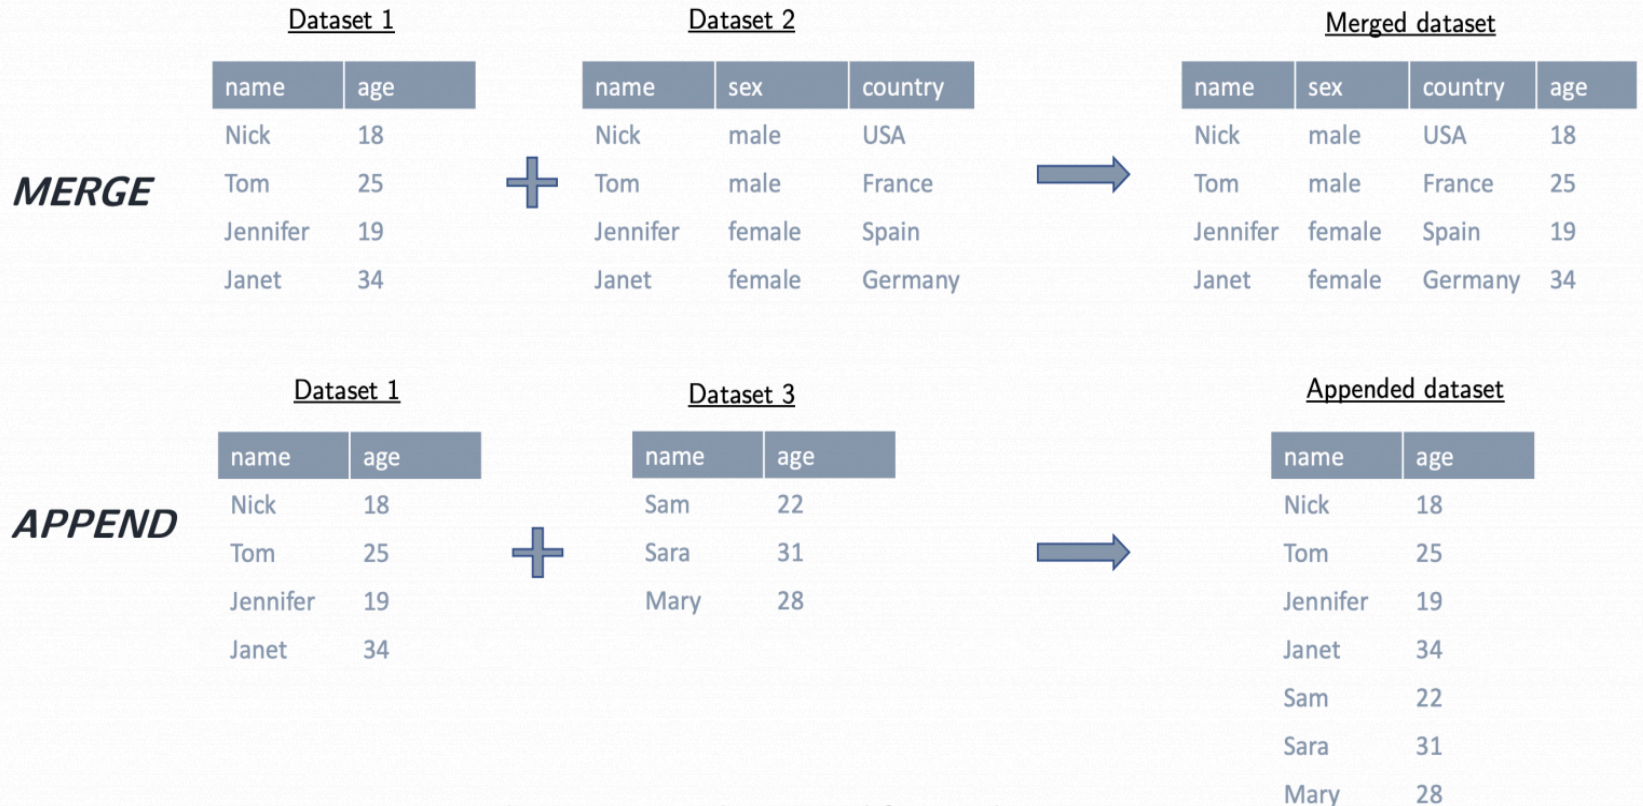
Write me code that will tells me the average birth year of characters by homeworld of the human characters in the `starwars` dataset.

Combining data frames

The final set of dplyr "goodies" are the family of **append** and **join** operations. However, these are important enough that I want to go over some concepts in a bit more depth...

- We will encounter and practice these many more times as the course progresses.
- Imagine you have two data frames, `df1` and `df2`, that you want to combine.
 - You can **append** or **bind**: stack the datasets on top of each other and match up the columns using `bind_rows()`
 - You can **merge** or **join**: match the rows based on a common identifier using `left_join()`, `inner_join()`, etc.
- The appropriate choice depends on the task you are trying to accomplish
 - Are you trying to add new observations or new variables?

Visualize the difference



Source: www.peretaberner.eu and @PereATaberner

Appending

- One way to append in the `tidyverse` is with `bind_rows()`
 - Base R has `rbind()`, which requires column names to match
 - `data.table` has `rbindlist()`, which requires column names to match unless you specify `fill`

```
df1 <- data.frame(x = 1:3, y = 4:6)
df2 <- data.frame(x = 1:4, y = 10:13, z=letters[1:4])
```

```
## Append df2 to df1
bind_rows(df1, df2)
```

```
##   x  y    z
## 1 1  4 <NA>
## 2 2  5 <NA>
## 3 3  6 <NA>
## 4 1 10    a
## 5 2 11    b
## 6 3 12    c
## 7 4 13    d
```

Joins

One of the mainstays of the dplyr package is merging data with the family [join operations](#).

- `inner_join(df1, df2)`
- `left_join(df1, df2)`
- `right_join(df1, df2)`
- `full_join(df1, df2)`
- `semi_join(df1, df2)`
- `anti_join(df1, df2)`

Joins are how you get **Relational Database Management** (RDBM) to work in R.

(See visual depictions of the different join operations [here](#).)

Joins (cont.)

Datasets to merge:

name	country		name	age
Nick	USA	↔	Nick	18
Tom	France	↔	Tom	25
Sara	France		Jennifer	19

Outputs:

inner_join()

name	country	age
Nick	USA	18
Tom	France	25

full_join()

name	country	age
Nick	USA	18
Tom	France	25
Sara	France	
Jennifer		19

left_join()

name	country	age
Nick	USA	18
Tom	France	25
Sara	France	

right_join()

name	country	age
Nick	USA	18
Tom	France	25
Jennifer		19

semi_join()

name	country
Nick	USA
Tom	France

anti_join()

name	country
Sara	France

Source: www.peretaberner.eu and @PereATaberner

Relational Database Management with R

- Remember relational database management?
- Each dataframe has a unique identifier (a "key") that links it to other dataframes.
- All the dataframes have the keys in common, so you can match them up
- Let's get a less abstract example using flights

nycflights13 data

The `flights` data frame contains information flights that departed from NYC in 2013.

- All flight information is stored in the `flights` data frame.
- Information about the planes (like year built) in the `planes` data frame.

```
## # A tibble: 6 × 6                                ## # A tibble: 6 × 4
##   flight tailnum  year month   day dep_time  ##   tailnum  year manufacturer    model
##   <int> <chr>    <int> <int> <int>   <int>  ##   <chr>    <int> <chr>         <chr>
## 1   1545 N14228   2013     1     1     517  ## 1 N10156   2004 EMBRAER        EMB-145XR
## 2   1714 N24211   2013     1     1     533  ## 2 N102UW   1998 AIRBUS INDUSTRIE A320-214
## 3   1141 N619AA   2013     1     1     542  ## 3 N103US   1999 AIRBUS INDUSTRIE A320-214
## 4    725 N804JB   2013     1     1     544  ## 4 N104UW   1999 AIRBUS INDUSTRIE A320-214
## 5    461 N668DN   2013     1     1     554  ## 5 N10575   2002 EMBRAER        EMB-145LR
## 6   1696 N39463   2013     1     1     554  ## 6 N105UW   1999 AIRBUS INDUSTRIE A320-214
```

Joins (cont.)

Let's perform a **left join** on the flights and planes datasets.

- *Note:* I'm going subset columns after the join, but only to keep text on the slide.

Joins (cont.)

Let's perform a **left join** on the flights and planes datasets.

- *Note:* I'm going subset columns after the join, but only to keep text on the slide.

```
left_join(flights, planes) %>%  
  select(year, month, day, dep_time, arr_time, carrier, flight, tailnum, type, model)
```

```
## Joining with by = join_by(year, tailnum)
```

```
## # A tibble: 336,776 × 10
```

```
##   year month   day dep_time arr_time carrier flight tailnum type  model  
##   <int> <int> <int>   <int>   <int> <chr>   <int> <chr>   <chr> <chr>  
## 1  2013     1     1     517     830 UA      1545 N14228 <NA> <NA>  
## 2  2013     1     1     533     850 UA      1714 N24211 <NA> <NA>  
## 3  2013     1     1     542     923 AA      1141 N619AA <NA> <NA>  
## 4  2013     1     1     544    1004 B6       725 N804JB <NA> <NA>  
## 5  2013     1     1     554     812 DL       461 N668DN <NA> <NA>  
## 6  2013     1     1     554     740 UA      1696 N39463 <NA> <NA>  
## 7  2013     1     1     555     913 B6       507 N516JB <NA> <NA>  
## 8  2013     1     1     557     709 EV      5708 N829AS <NA> <NA>  
## 9  2013     1     1     557     838 B6        79 N593JB <NA> <NA>  
## 10 2013     1     1     558     753 AA       301 N3ALAA <NA> <NA>
```

```
## # i 336,766 more rows
```

Joins (cont.)

(continued from previous slide)

Note that dplyr made a reasonable guess about which columns to join on (i.e. columns that share the same name). It also told us its choices:

```
## Joining, by = c("year", "tailnum")
```

However, there's a problem here: the variable "year" does not have a consistent meaning across our joining datasets!

- In one it refers to the *year of flight*, in the other it refers to *year of construction*.

Joins (cont.)

(continued from previous slide)

Note that dplyr made a reasonable guess about which columns to join on (i.e. columns that share the same name). It also told us its choices:

```
## Joining, by = c("year", "tailnum")
```

However, there's a problem here: the variable "year" does not have a consistent meaning across our joining datasets!

- In one it refers to the *year of flight*, in the other it refers to *year of construction*.

Luckily, there's an easy way to avoid this problem.

- See if you can figure it out before turning to the next slide.
- Get help with `?dplyr::join`

Joins (cont.)

(continued from previous slide)

You just need to be more explicit in your join call by using the `by =` argument.

- You can also rename any ambiguous columns to avoid confusion.

```
left_join(
  flights,
  planes %>% rename(year_built = year), ## Not necessary w/ below line, but helpful
  by = "tailnum" ## Be specific about the joining column
) %>%
  select(year, month, day, dep_time, arr_time, carrier, flight, tailnum, year_built, type, model)
  head(3) ## Just to save vertical space on the slide
```

```
## # A tibble: 3 × 11
##   year month   day dep_time arr_time carrier flight tailnum year_built type
##   <int> <int> <int>   <int>   <int> <chr>   <int> <chr>      <int> <chr>
## 1  2013     1     1     517     830 UA      1545 N14228     1999 Fixed w...
## 2  2013     1     1     533     850 UA      1714 N24211     1998 Fixed w...
## 3  2013     1     1     542     923 AA      1141 N619AA     1990 Fixed w...
## # i 1 more variable: model <chr>
```

Joins (cont.)

(continued from previous slide)

Last thing I'll mention for now; note what happens if we again specify the join column... but don't rename the ambiguous "year" column in at least one of the given data frames.

```
left_join(
  flights,
  planes, ## Not renaming "year" to "year_built" this time
  by = "tailnum"
) %>%
select(contains("year"), month, day, dep_time, arr_time, carrier, flight, tailnum, type, model)
head(3)
```

```
## # A tibble: 3 × 11
##   year.x year.y month   day dep_time arr_time carrier flight tailnum type  model
##   <int>  <int> <int> <int>   <int>   <int> <chr>   <int> <chr>  <chr> <chr>
## 1  2013   1999     1     1     517     830 UA       1545 N14228 Fixe... 737-...
## 2  2013   1998     1     1     533     850 UA       1714 N24211 Fixe... 737-...
## 3  2013   1990     1     1     542     923 AA       1141 N619AA Fixe... 757-...
```

Joins (cont.)

(continued from previous slide)

Last thing I'll mention for now; note what happens if we again specify the join column... but don't rename the ambiguous "year" column in at least one of the given data frames.

```
left_join(
  flights,
  planes, ## Not renaming "year" to "year_built" this time
  by = "tailnum"
) %>%
select(contains("year"), month, day, dep_time, arr_time, carrier, flight, tailnum, type, model)
head(3)
```

```
## # A tibble: 3 × 11
##   year.x year.y month   day dep_time arr_time carrier flight tailnum type  model
##   <int>  <int> <int> <int>   <int>   <int> <chr>   <int> <chr>  <chr> <chr>
## 1  2013   1999     1     1     517     830 UA       1545 N14228 Fixe... 737-...
## 2  2013   1998     1     1     533     850 UA       1714 N24211 Fixe... 737-...
## 3  2013   1990     1     1     542     923 AA       1141 N619AA Fixe... 757-...
```

Make sure you know what "year.x" and "year.y" are. Again, it pays to be specific.

tidyr

Key tidyr verbs

1. `pivot_longer`: Pivot wide data into long format.
2. `pivot_wider`: Pivot long data into wide format.
3. `separate`, `unite`, `fill`, `expand`, `nest`, `unnest`: Various other data tidying operations.
 - There are many utilities in the `tidyr` package that help you clean and wrangle data.
 - But they are best learned through experience

Key tidyr verbs

1. `pivot_longer`: Pivot wide data into long format.
2. `pivot_wider`: Pivot long data into wide format.
3. `separate`, `unite`, `fill`, `expand`, `nest`, `unnest`: Various other data tidying operations.
 - There are many utilities in the `tidyr` package that help you clean and wrangle data.
 - But they are best learned through experience

Let's practice these verbs together in class.

- Side question: Which of `pivot_longer` vs `pivot_wider` produces "tidy" data?

1) tidyr::pivot_longer

```
stocks = data.frame( ## Could use "tibble" instead of "data.frame" if you prefer
  time = as.Date('2009-01-01') + 0:1,
  X = rnorm(2, 0, 1), Y = rnorm(2, 0, 2), Z = rnorm(2, 0, 4))
stocks
```

```
##           time           X           Y           Z
## 1 2009-01-01  0.4139186 -0.3254475  2.087752
## 2 2009-01-02 -1.2610702 -3.8178951 -3.455760
```

```
tidy_stocks = stocks %>% pivot_longer(-time, names_to="stock", values_to="price")
tidy_stocks
```

```
## # A tibble: 6 × 3
##   time      stock price
##   <date>   <chr> <dbl>
## 1 2009-01-01 X      0.414
## 2 2009-01-01 Y     -0.325
## 3 2009-01-01 Z      2.09
## 4 2009-01-02 X     -1.26
## 5 2009-01-02 Y     -3.82
## 6 2009-01-02 Z     -3.46
```

2) tidyr::pivot_wider

```
tidy_stocks %>% pivot_wider(names_from=stock, values_from=price)
```

```
## # A tibble: 2 × 4
##   time          X      Y      Z
##   <date>      <dbl> <dbl> <dbl>
## 1 2009-01-01 0.0231 -2.08 -2.22
## 2 2009-01-02 1.25   -3.45  6.01
```

```
tidy_stocks %>% pivot_wider(names_from=time, values_from=price)
```

```
## # A tibble: 3 × 3
##   stock 2009-01-01 2009-01-02
##   <chr>      <dbl>      <dbl>
## 1 X          0.0231         1.25
## 2 Y          -2.08        -3.45
## 3 Z          -2.22         6.01
```

2) tidyr::pivot_wider

```
tidy_stocks %>% pivot_wider(names_from=stock, values_from=price)
```

```
## # A tibble: 2 × 4  
##   time          X      Y      Z  
##   <date>      <dbl> <dbl> <dbl>  
## 1 2009-01-01 0.0231 -2.08 -2.22  
## 2 2009-01-02 1.25   -3.45  6.01
```

```
tidy_stocks %>% pivot_wider(names_from=time, values_from=price)
```

```
## # A tibble: 3 × 3  
##   stock 2009-01-01 2009-01-02  
##   <chr>      <dbl>      <dbl>  
## 1 X          0.0231         1.25  
## 2 Y          -2.08        -3.45  
## 3 Z          -2.22         6.01
```

Note that the second example — which has combined different pivoting arguments — has effectively transposed the data.

2) tidyr::pivot_longer with prefix

Let's pivot the pre-loaded billboard data: showing weekly rankings of top 100 in the year 2000

```
head(billboard)
```

```
## # A tibble: 6 × 79
##   artist      track date.entered  wk1  wk2  wk3  wk4  wk5  wk6  wk7  wk8
##   <chr>      <chr> <date>      <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 2 Pac      Baby... 2000-02-26      87   82   72   77   87   94   99   NA
## 2 2Ge+her    The ... 2000-09-02      91   87   92   NA   NA   NA   NA   NA
## 3 3 Doors Do... Kryp... 2000-04-08      81   70   68   67   66   57   54   53
## 4 3 Doors Do... Loser 2000-10-21      76   76   72   69   67   65   55   59
## 5 504 Boyz    Wobb... 2000-04-15      57   34   25   17   17   31   36   49
## 6 98^0       Give... 2000-08-19      51   39   34   26   26   19    2    2
## # i 68 more variables: wk9 <dbl>, wk10 <dbl>, wk11 <dbl>, wk12 <dbl>,
## #   wk13 <dbl>, wk14 <dbl>, wk15 <dbl>, wk16 <dbl>, wk17 <dbl>, wk18 <dbl>,
## #   wk19 <dbl>, wk20 <dbl>, wk21 <dbl>, wk22 <dbl>, wk23 <dbl>, wk24 <dbl>,
## #   wk25 <dbl>, wk26 <dbl>, wk27 <dbl>, wk28 <dbl>, wk29 <dbl>, wk30 <dbl>,
## #   wk31 <dbl>, wk32 <dbl>, wk33 <dbl>, wk34 <dbl>, wk35 <dbl>, wk36 <dbl>,
## #   wk37 <dbl>, wk38 <dbl>, wk39 <dbl>, wk40 <dbl>, wk41 <dbl>, wk42 <dbl>,
## #   wk43 <dbl>, wk44 <dbl>, wk45 <dbl>, wk46 <dbl>, wk47 <dbl>, wk48 <dbl>, ...
```

2) tidyr::pivot_longer with prefix *cont.*

Wait, why is there 'wk' in the 'week' column?

```
billboard %>%  
  pivot_longer(cols=starts_with('wk'), names_to="week",  
    values_to="rank") %>%  
  head()
```

```
## # A tibble: 6 × 5  
##   artist track          date.entered week    rank  
##   <chr> <chr>          <date>      <chr> <dbl>  
## 1 2 Pac Baby Don't Cry (Keep ... 2000-02-26 wk1      87  
## 2 2 Pac Baby Don't Cry (Keep ... 2000-02-26 wk2      82  
## 3 2 Pac Baby Don't Cry (Keep ... 2000-02-26 wk3      72  
## 4 2 Pac Baby Don't Cry (Keep ... 2000-02-26 wk4      77  
## 5 2 Pac Baby Don't Cry (Keep ... 2000-02-26 wk5      87  
## 6 2 Pac Baby Don't Cry (Keep ... 2000-02-26 wk6      94
```

2) tidyr::pivot_longer with prefix cont.

That fixed it.

```
billboard %>%  
  pivot_longer(cols=starts_with('wk'), names_to="week",  
    values_to="rank",names_prefix='wk') %>%  
  mutate(week=as.numeric(week)) %>% # Make week a numeric variable  
  head()
```

```
## # A tibble: 6 × 5  
##   artist track          date.entered  week  rank  
##   <chr>  <chr>          <date>      <dbl> <dbl>  
## 1 2 Pac  Baby Don't Cry (Keep ... 2000-02-26      1     87  
## 2 2 Pac  Baby Don't Cry (Keep ... 2000-02-26      2     82  
## 3 2 Pac  Baby Don't Cry (Keep ... 2000-02-26      3     72  
## 4 2 Pac  Baby Don't Cry (Keep ... 2000-02-26      4     77  
## 5 2 Pac  Baby Don't Cry (Keep ... 2000-02-26      5     87  
## 6 2 Pac  Baby Don't Cry (Keep ... 2000-02-26      6     94
```

Aside: Remembering the `pivot_*` syntax

There's a long-running joke about no-one being able to remember Stata's "reshape" command. ([Exhibit A](#).)

It's easy to see this happening with the `pivot_*` functions too. Remember the documentation is your friend!

```
?pivot_longer
```

And GitHub CoPilot, ChatGPT and other AI tools are also your friends if you use precise language about what you want the AI tool to do and you try their suggestions carefully.¹

¹ Back in my day we had to scour StackOverflow for hours to find the right answer. And we liked it!

Other tidyr goodies

- `separate`: Split a single column into multiple columns.
 - `separate(df, col, into = c("A", "B"), sep = "-")` will split `col` into columns `A` and `B` at the `-` separator.
- `unite`: Combine multiple columns into a single column.
 - `unite(df, col, A, B, sep = "-")` combines columns `A` and `B` into column `col` with `-` as the separator.
- `fill`: Fill in missing values with the last non-missing value.
 - `fill(df, starts_with("X"))` will fill in all columns that start with "X".
- `drop_na`: Drop rows with missing values.
- `expand`: Create a complete set of combinations from a set of factors.
- `nest` and `unnest`: Combine columns into lists within a single cell or split a column of lists into separate rows.
 - Try with the `starwars` data frame: `unnest(starwars, films, names_sep='')`

Summary

Key verbs

dplyr

1. `filter`
2. `arrange`
3. `select`
4. `mutate`
5. `summarise`

tidyr

1. `pivot_longer`
2. `pivot_wider`

Key verbs

dplyr

1. `filter`
2. `arrange`
3. `select`
4. `mutate`
5. `summarise`

tidyr

1. `pivot_longer`
2. `pivot_wider`

Other useful items include: pipes (`%>%`), grouping (`group_by`), joining functions (`left_join`, `inner_join`, etc.).

Datacamp!

- Navigate to datacamp (link on class materials ReadMe)

Next lecture: Scraping data!
